

# Introduction to Aerial Robotics

## Lecture 5

Shaojie Shen  
Associate Professor  
Dept. of ECE, HKUST



4 March 2024

# Course Outline

- Dynamics, Planning & Control

- Vision

- Estimation

Week	Lecture Date Tue 13:30-16:20	Topic	Assignment (due <b>23:59</b> on <b>Friday</b> of the corresponding week)	Lab Wed 18:00-20:50
1	2/4	Introduction		No Lab
2	2/11	Rigid Body Transformation Quaternion Quadrotor Modeling		No Lab
3	2/18	Control Basics Quadrotor Control Trajectory Generation	Project 1 Phase 1 Out	No Lab
4	2/25	Trajectory Generation Path Planning	Project 1 Phase 1 Due Project 1 Phase 2 Out	No Lab
5	3/4	Camera Modeling & Calibration Feature Detection & Matching	Project 1 Phase 2 Due Project 1 Phase 3 Out	No Lab
6	3/11	Midterm	Project 1 Phase 3 Due Project 1 Phase 4 Out	Lab Tutorial 1: Robot Assembly
7	3/18	Multi-View Geometry Pose Estimation	Project 1 Phase 3 Due Project 2 Phase 1 Out	Lab Tutorial 2: Prepare P1P4
8	3/25	Optical Flow Dense Stereo	Project 2 Phase 1 Due Project 2 Phase 2 Out	Free Lab Time
9	4/1	Midterm Break (No class)		No Lab
10	4/8	Probability Basics Bayesian Inferencing Kalman Filter	Project 2 Phase 2 Due Project 3 Phase 1 Out	Free Lab Time
11	4/15	Extended Kalman Filter Augmented State EKF Particle Filter	Project 1 Phase 4 Due Project 3 Phase 2 Out	Free Lab Time
12	4/22	SLAM	Project 3 Phase 3 Out Project 3 Phase 1 Due	Lab Tutorial 3: Prepare P3P3
13	4/29	x	Project 3 Phase 2 Due	Free Lab Time
14	5/6	x	Project 3 Phase 3 Due	Free Lab Time

# Outline


- Camera Modeling
- Robot Vision Pipeline
- Point Feature Detection & Matching

# Camera Modeling

# Goal: Fly Like Birds



# Robot Sees with Cameras



# Cameras


- Monocular

- Simplest setup
- Depth unknown



- Stereo

- Able to compute depth
- Depth accuracy affected by baseline, resolution, and calibration




- Multi-Camera

- Overlapping / Non-overlapping field-of-view



Omnidirectional multi-camera system  
Ladybug



Relative position and posture of each camera



# Cameras

- RGB-D Sensor
  - Great depth
  - Does not work outdoors



- Omnidirectional camera
  - 360 capture
  - Strong distortion







# Cameras


- Sensor

- Lens




# Thin Lens

- A lens with a thickness that is negligible compared to the radii of curvature of the lens surfaces



# Thin Lens


- Rays parallel to the optical axis meet focus after leaving the lens
- Rays through center of the lens do not change direction



$$\frac{1}{f} = \frac{1}{a} + \frac{1}{b}$$


# Thin Lens

- De-focus if  $\frac{1}{f} \neq \frac{1}{a} + \frac{1}{b}$





# Pin-hole Camera Model




WIKIPEDIA  
The Free Encyclopedia

# Pin-hole Camera Model

- If we replace  $b$  with  $f$  and include a minus because the object image is upside down ( $Z = a, f = b$ )

$$-y = -f \frac{Y}{Z}$$






# Pin-hole Camera Model

- Assume that the image plane is in front of the lens:


$$-y = f \frac{Y}{Z}$$



# Effect of $f = b$ ?

- Theoretically ,we expect an offset in the x and y coordinates caused by the error ( $f - b$ )
- If the object is on focus:  $\frac{b-f}{f} = \frac{b}{Z}$ 
  - Relative error depends on the ratio of focus length to depth.
  - This matters if we actually use the focus length from camera specs
  - In practice, we use a procedure called **calibration** to obtain  $f$  that best satisfies:  $y = f \frac{Y}{Z}$

# Camera Coordinate System




# Perspective Projection

- Optical axis is the z-axis
- The image plane  $(u, v)$  is perpendicular to the optical axis
- Intersection of the image plane and the optical axis is the image center  $(u_0, v_0)$
- $f$  is the distance of the image plane from the origin (in pixels)
- Formulation:

$$- u = \frac{fX_c}{Z_c} + u_0$$


$$- v = \frac{fY_c}{Z_c} + v_0$$



# Perspective Projection

- Matrix form:


$$\lambda \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} X_c \\ Y_c \\ Z_c \\ 1 \end{pmatrix}$$




# Perspective Projection

- From camera to world:

$$\begin{pmatrix} X_c \\ Y_c \\ Z_c \\ 1 \end{pmatrix} = \begin{pmatrix} R & t \\ 0 & 1 \end{pmatrix} \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$





# Pin-hole Camera Model






# Radial Distortions

- Wide angle lenses have radial distortions
  - Straight lines become curves
- Formulation:
  - $r^2 = u^2 + v^2$
  - $u^{dist} = u(1 + k_1r + k_2r^2 + k_3r^3 + \dots)$
  - $v^{dist} = v(1 + k_1r + k_2r^2 + k_3r^3 + \dots)$



# Camera Calibration

- Requires:
  - Calibration object
- Obtains:
  - Intrinsic parameters
  - Distortion parameters
  - Poses of cameras with respect to the calibration object




$$\lambda \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix} (R \quad t) \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix} = P \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$

The diagram illustrates the camera calibration equation with color-coded components and their roles:

- Pixel Values** (blue box): Points to the vector  $\begin{pmatrix} u \\ v \\ 1 \end{pmatrix}$ .
- Intrinsic Calibration Matrix** (red box): Points to the matrix  $\begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix}$ .
- Camera Pose** (red box): Points to the matrix  $(R \quad t)$ .
- World Point** (green box): Points to the vector  $\begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$ .
- Known** (green box): Points to the projection matrix  $P$ .
- Measurement** (blue box): Points to the vector  $\begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$ .
- To be Estimated** (red box): Points to the projection matrix  $P$ .


# Camera Calibration

- Straight lines should be straight




# Robot Vision Pipeline

# Vision-based State Estimation Pipeline





# Feature Detection

- We cannot process the entire image directly
- Requirements:
  - Repeatability
  - Saliency
  - Locality
  - Compactness and efficiency
- Popular features
  - Corners (FAST, Harris, ...)
  - Blob (SIFT, SURF, ...)
  - Line (Canny, ...)




# Feature Matching

- Match features in different images
  - Across multiple cameras
  - Across time
- Common methods:
  - Descriptor matching
  - Optical flow






# Pose Estimation



$$\lambda \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} R & t \end{pmatrix} \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix} = P \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$

Pixel Values      Intrinsic Calibration Matrix      Camera Pose      World Point

# Visual Odometry



Known

Measurement

To be Estimated


$$\lambda \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} R & t \end{pmatrix} \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix} = P \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$

Pixel Values

Intrinsic Calibration Matrix

Camera (Incremental) Pose

# SLAM



$$\lambda \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} = \begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix} (R \quad t) \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix} = P \begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$$

Pixel Values (blue arrow pointing to  $\begin{pmatrix} u \\ v \\ 1 \end{pmatrix}$ )

Intrinsic Calibration Matrix (green arrow pointing to  $\begin{pmatrix} f & 0 & u_o \\ 0 & f & v_o \\ 0 & 0 & 1 \end{pmatrix}$ )

Camera Pose (red arrow pointing to  $(R \quad t)$ )

World Point (red arrow pointing to  $\begin{pmatrix} X_w \\ Y_w \\ Z_w \\ 1 \end{pmatrix}$ )

# Point Feature Detection & Matching

# Image matching



by [Diva Sian](#)



by [swashford](#)





# Harder case



by [Diva Sian](#)





by [scgbt](#)

# Even harder case




***“How the Afghan Girl was Identified by Her Iris Patterns”*** Read the [story](#)






# Harder still?



NASA Mars Rover images


# Answer below (look for tiny colored squares...)



NASA Mars Rover images  
with SIFT feature matches  
Figure by Noah Snavely






# Image Matching





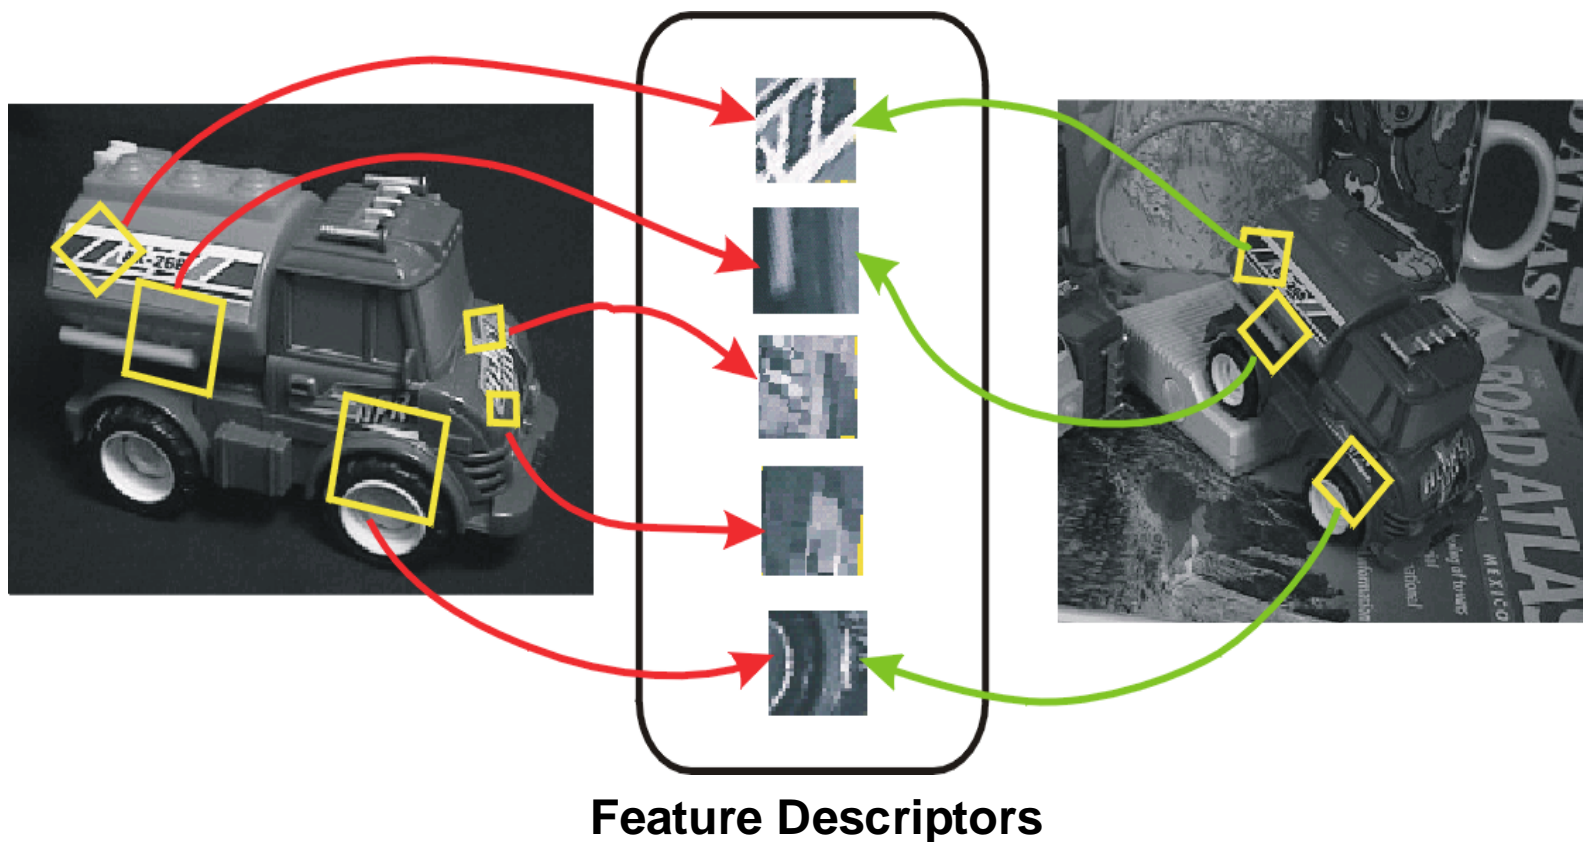


# Image Matching



# Invariant local features

- Find features that are invariant to transformations
  - geometric invariance: translation, rotation, scale
  - photometric invariance: brightness, exposure, ...




# Advantages of local features

- Locality
  - Features are local, so robust to occlusion and clutter
- Distinctiveness
  - Can differentiate a large database of objects
- Quantity
  - Hundreds or thousands in a single image
- Efficiency
  - Real-time performance achievable
- Generality
  - Exploit different types of features in different situations

# More motivation...

- Feature points are used for:
  - Image alignment (e.g., mosaics)
  - 3D reconstruction
  - Motion tracking
  - Object recognition
  - Indexing and database retrieval
  - Robot navigation
  - ... other

# Image Mosaics

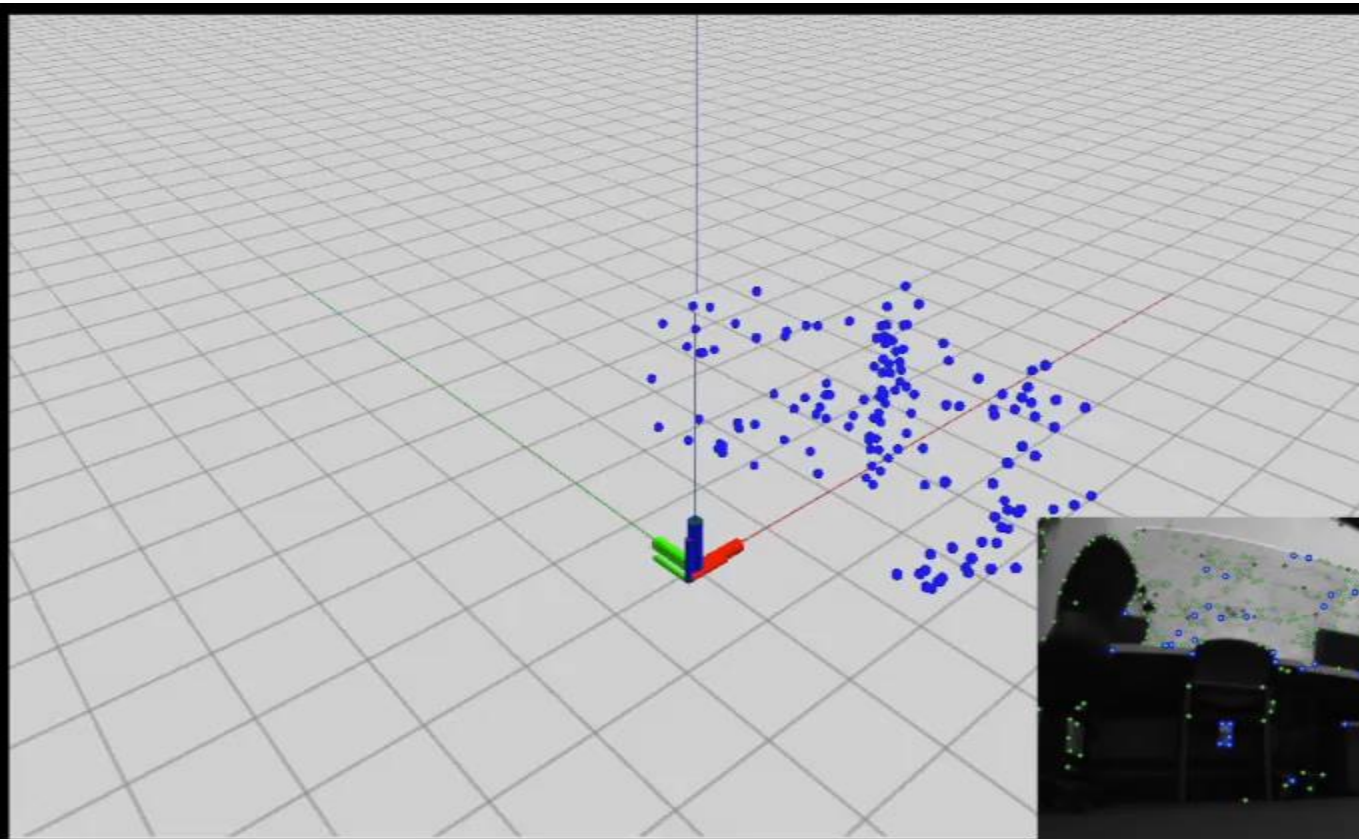





# Image Mosaics



# Motion Tracking



Large axes: Vision-only pose  
Small axes & arrow: Vision-IMU pose & velocity  
Blue line: Vision-only trajectory  
Yellow line: Vision-IMU trajectory  
Red & blue dots: 3D features




# Want uniqueness

- Look for image regions that are unusual
  - Lead to unambiguous matches in other images
- How to define “unusual”?

# Local measures of uniqueness

Suppose we only consider a small window of pixels


- What defines whether a feature is a good or bad candidate?




Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

# Feature detection


- Local measure of feature uniqueness
  - How does the window change when you shift it?
  - Shifting the window in *any direction* causes a *big change*



“flat” region:  
no change in all  
directions





“edge”:  
no change along the  
edge direction



“corner”:  
significant change in  
all directions


Slide adapted from Darya Frolova, Denis Simakov, Weizmann Institute.

# Feature detection



# Feature detection: the math

- Consider shifting the window **W** by  $(u,v)$ 
  - How do the pixels in **W** change?
  - Compare each pixel before and after by summing up the squared differences (SSD)
  - This defines an SSD “error” of  $E(u,v)$ :









$$E(u, v) = \sum_{(x,y) \in W} [I(x + u, y + v) - I(x, y)]^2$$





(a)





# Small motion assumption

- Taylor Series expansion of I:

$$I(x + u, y + v) = I(x, y) + \frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v + \text{higher order terms}$$


- If the motion (u,v) is small, then first order approximation is good

$$\begin{aligned} I(x + u, y + v) &\approx I(x, y) + \frac{\partial I}{\partial x} u + \frac{\partial I}{\partial y} v \\ &\approx I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix} \end{aligned}$$


$$\text{shorthand: } I_x = \frac{\partial I}{\partial x}$$

- Plugging this into the formula on the previous slide...


# Image Gradients



$I$




$$I_x = I(x + 1, y) - I(x - 1, y)$$



$$I_y = I(x, y + 1) - I(x, y - 1)$$

# Feature detection: the math

- Consider shifting the window  $\mathbf{W}$  by  $(u, v)$ 
  - How do the pixels in  $\mathbf{W}$  change?
  - Compare each pixel before and after by summing up the squared differences
  - This defines an “error” of  $E(u, v)$ :




$$\begin{aligned}
 E(u, v) &= \sum_{(x, y) \in W} [I(x + u, y + v) - I(x, y)]^2 \\
 &\approx \sum_{(x, y) \in W} \left[ I(x, y) + \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} - I(x, y) \right]^2 \\
 &\approx \sum_{(x, y) \in W} \left[ \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} \right]^2
 \end{aligned}$$

# Feature detection: the math

- This can be rewritten:

$$E(u, v) = \sum_{(x,y) \in W} [u \quad v] \underbrace{\begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$



- For the example above
  - You can move the center of the green window to anywhere on the blue unit circle
  - Which directions will result in the largest and smallest  $E$  values?
  - We can find these directions by looking at the eigenvectors of  $H$

# Quick eigenvalue/eigenvector review

- The **eigenvectors** of a matrix **A** are the vectors **x** that satisfy:

$$Ax = \lambda x$$

- The scalar  $\lambda$  is the **eigenvalue** corresponding to **x**

- The eigenvalues are found by solving:

$$\det(A - \lambda I) = 0$$

- In our case, **A** = **H** is a 2x2 matrix, so we have

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} = 0$$

- The solution:

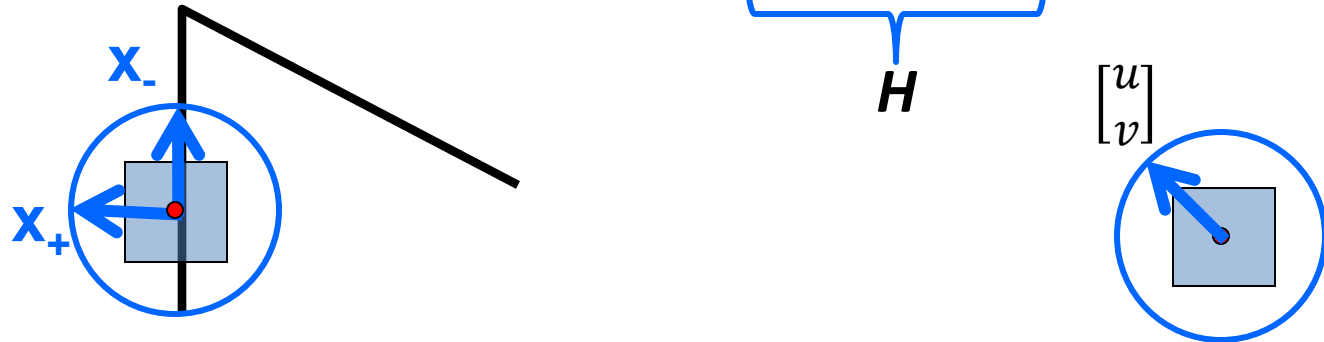
$$\lambda_{\pm} = \frac{1}{2} \left[ (h_{11} + h_{22}) \pm \sqrt{4h_{12}h_{21} + (h_{11} - h_{22})^2} \right]$$

- Once you know  $\lambda$ , you find **x** by solving

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{21} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

# Feature detection: the math

- This can be rewritten:

$$E(u, v) = \sum_{(x,y) \in W} [u \quad v] \underbrace{\begin{bmatrix} I_x^2 & I_x I_y \\ I_y I_x & I_y^2 \end{bmatrix}}_H \begin{bmatrix} u \\ v \end{bmatrix}$$



- Eigenvalues and eigenvectors of  $H$ 
  - Define shifts with the smallest and largest change (E value)
  - $x_+$  = direction of **largest** increase in E.
  - $\lambda_+$  = amount of increase in direction  $x_+$
  - $x_-$  = direction of **smallest** increase in E.
  - $\lambda_-$  = amount of increase in direction  $x_+$

$$Hx_+ = \lambda_+ x_+$$

$$Hx_- = \lambda_- x_-$$


# Feature detection: the math

- How are  $\lambda_+$ ,  $\mathbf{x}_+$ ,  $\lambda_-$ , and  $\mathbf{x}_-$  relevant for feature detection?
  - What's our feature scoring function?




# Feature detection: the math


- How are  $\lambda_+$ ,  $\mathbf{x}_+$ ,  $\lambda_-$ , and  $\mathbf{x}_-$  relevant for feature detection?
  - What's our feature scoring function?
- Want  $E(u,v)$  to be **large** for small shifts in **all** directions
  - the *minimum* of  $E(u,v)$  should be large, over all unit vectors  $[u \ v]$
  - this minimum is given by the smaller eigenvalue ( $\lambda_-$ ) of  $\mathbf{H}$



$I$



$\lambda_+$




$\lambda_-$




# Feature detection summary


- Here's what you do
  - Compute the gradient at each point in the image
  - Create the ***H*** matrix from the entries in the gradient
  - Compute the eigenvalues.
  - Find points with large response ( $\lambda_- > \text{threshold}$ )
  - Choose those points where  $\lambda_-$  is a local maximum as features



$I$




$\lambda_+$



$\lambda_-$

# Feature detection summary

- Here's what you do
  - Compute the gradient at each point in the image
  - Create the ***H*** matrix from the entries in the gradient
  - Compute the eigenvalues.
  - Find points with large response ( $\lambda_- > \text{threshold}$ )
  - Choose those points where  $\lambda_-$  is a local maximum as features



$\lambda_-$


# The Harris operator

- $\lambda_-$  is a variant of the “Harris operator” for feature detection

$$\begin{aligned} f &= \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} \\ &= \frac{\text{determinant}(H)}{\text{trace}(H)} \end{aligned}$$

- The *trace* is the sum of the diagonals, i.e.,  $\text{trace}(H) = h_{11} + h_{22}$
- Very similar to  $\lambda_-$  but less expensive (no square root)
- Called the “Harris Corner Detector” or “Harris Operator”
- Lots of other detectors, this is one of the most popular

# The Harris operator




# Harris detector example





# f value (red high, blue low)



# Threshold ( $f > \text{value}$ )



# Find local maxima of $f$





# Harris features (in red)




# Invariance

- Suppose you **rotate** the image by some angle
  - Will you still pick up the same features?
- What if you change the brightness?
- Scale?

# Scale invariant detection

- Suppose you're looking for corners




- Key idea: find scale that gives local maximum of  $f$ 
  - $f$  is a local maximum in both position and scale

# Automatic scale selection





Lindeberg et al., 1996





$$f(I_{i_1 \dots i_m}(x, \sigma))$$

# Automatic scale selection



$$f(I_{i_1 \dots i_m}(x, \sigma))$$


# Automatic scale selection



$$f(I_{i_1 \dots i_m}(x, \sigma))$$




# Automatic scale selection



$$f(I_{i_1 \dots i_m}(x, \sigma))$$




# Automatic scale selection




$$f(I_{i_1...i_m}(x, \sigma))$$

# Automatic scale selection




$$f(I_{i_1 \dots i_m}(x, \sigma))$$

# Automatic scale selection



$$f(I_{i_1 \dots i_m}(x, \sigma))$$



$$f(I_{i_1 \dots i_m}(x', \sigma'))$$

# Automatic scale selection

- Normalize: rescale to fixed size






# Feature descriptors

We know how to detect good points


Next question: **How to match them?**



# Feature descriptors

We know how to detect good points

Next question: **How to match them?**



Lots of possibilities (this is a popular research area)

- Simple option: match square windows around the point
- Better approach: SIFT
  - David Lowe, UBC <http://www.cs.ubc.ca/~lowe/keypoints/>

# Invariance

Suppose we are comparing two images  $I_1$  and  $I_2$

- $I_2$  may be a transformed version of  $I_1$
- What kinds of transformations are we likely to encounter in practice?





# Invariance

- Suppose we are comparing two images  $I_1$  and  $I_2$ 
  - $I_2$  may be a transformed version of  $I_1$
  - What kinds of transformations are we likely to encounter in practice?
- We'd like to find the same features regardless of the transformation
  - This is called transformational ***invariance***
  - Most feature methods are designed to be invariant to
    - Translation, 2D rotation, scale
  - They can usually also handle
    - Limited 3D rotations (SIFT works up to about 60 degrees)
    - Limited affine transformations (some are fully affine invariant)
    - Limited illumination/contrast changes



# How to achieve invariance

Need both of the following:

## 1. Make sure your detector is invariant

- Harris is invariant to translation and rotation
- Scale is trickier
  - common approach is to detect features at many scales using a Gaussian pyramid (e.g., MOPS)
  - More sophisticated methods find “the best scale” to represent each feature (e.g., SIFT)

## 2. Design an invariant feature *descriptor*

- A descriptor captures the information in a region around the detected feature point
- The simplest descriptor: a square window of pixels
  - What’s this invariant to?
- Let’s look at some better approaches...

# Rotation invariance for feature descriptors

Find dominant orientation of the image patch

- This is given by  $\mathbf{x}_+$ , the eigenvector of  $\mathbf{H}$  corresponding to  $\lambda_+$ 
  - $\lambda_+$  is the *larger* eigenvalue
- Rotate the patch according to this angle




Figure by Matthew Brown

# Multiscale Oriented Patches descriptor


Take 40x40 square window around detected feature

- Scale to 1/5 size (using prefiltering)
- Rotate to horizontal
- Sample 8x8 square window centered at feature
- Intensity normalize the window by subtracting the mean, dividing by the standard deviation in the window



Slide adapted from Matthew Brown

# Detections at multiple scales




*Figure 1. Multi-scale Oriented Patches (MOPS) extracted at five pyramid levels from one of the Matier images. The boxes show the feature orientation and the region from which the descriptor vector is sampled.*



# Scale Invariant Feature Transform


- Basic idea:
  - Take 16x16 square window around detected feature
  - Compute edge orientation (angle of the gradient -  $90^\circ$ ) for each pixel
  - Throw out weak edges (threshold gradient magnitude)
  - Create histogram of surviving edge orientations



Slide adapted from David Lowe

# SIFT descriptor

- Full version
  - Divide the 16x16 window into a 4x4 grid of cells (2x2 case shown below)
  - Compute an orientation histogram for each cell
  - 16 cells \* 8 orientations = 128 dimensional descriptor



Slide adapted from David Lowe





# Properties of SIFT

- Extraordinarily robust matching technique
  - Can handle changes in viewpoint
    - Up to about 60 degree out of plane rotation
  - Can handle significant changes in illumination
    - Sometimes even day vs. night (below)
  - Fast and efficient—can run in real time
  - Lots of code available
    - [http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known\\_implementations\\_of\\_SIFT](http://people.csail.mit.edu/albert/ladypack/wiki/index.php/Known_implementations_of_SIFT)


# Feature matching

Given a feature in  $I_1$ , how to find the best match in  $I_2$ ?


1. Define distance function that compares two descriptors
2. Test all the features in  $I_2$ , find the one with min distance

# Feature distance

- How to define the difference between two features  $f_1, f_2$ ?
  - Simple approach is  $SSD(f_1, f_2)$ 
    - sum of square differences between entries of the two descriptors
    - can give good scores to very ambiguous (bad) matches




$I_1$




$I_2$

# Feature distance

- How to define the difference between two features  $f_1, f_2$ ?
  - Better approach: ratio distance =  $\text{SSD}(f_1, f_2) / \text{SSD}(f_1, f_2')$ 
    - $f_2$  is best SSD match to  $f_1$  in  $I_2$
    - $f_2'$  is 2<sup>nd</sup> best SSD match to  $f_1$  in  $I_2$
    - gives small values for ambiguous matches




$I_1$



$I_2$

# Evaluating the results


How can we measure the performance of a feature matcher?



feature distance




# True/false positives



- The distance threshold affects performance
  - True positives = # of detected matches that are correct
    - Suppose we want to maximize these—how to choose threshold?
  - False positives = # of detected matches that are incorrect
    - Suppose we want to minimize these—how to choose threshold?

# Evaluating the results

How can we measure the performance of a feature matcher?






# Evaluating the results

- How can we measure the performance of a feature matcher?

## ROC curve ("Receiver Operator Characteristic")



## ROC Curves

- Generated by counting # current/incorrect matches, for different thresholds
- Want to maximize area under the curve (AUC)
- Useful for comparing different feature matching methods

# Lots of applications

- Features are used for:
  - Image alignment (e.g., mosaics)
  - 3D reconstruction
  - Motion tracking
  - Object recognition
  - Indexing and database retrieval
  - Robot navigation
  - ... other

# Logistics

- Project 1, phase 2 due this Friday (03/07)
- Project 1, phase 3 due next Friday (03/14)
- Lab tutorial next Wednesday (3/12)
  - 6:00pm-8:50pm @ Robotics Institute G/F, Flight Area
  - 2~3 students per group.
  - Try your best to find your groupmates, or we will randomly assign for you.
- Midterm next Tuesday (03/11)
  - 1:30pm-3:30pm @ Rm5506
  - Open book, open notes, close Internet
  - No communication with your classmates, honor code is strictly enforced
  - Covers lectures 1-4
  - 2 hours exam